


# Day 8: Image Classification

## Classification and the Tendril Problem

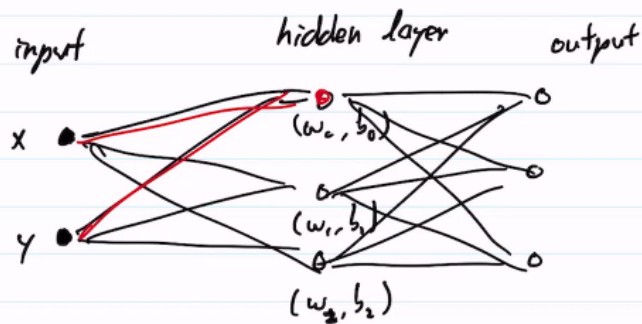


Example of Multidimensional Input/Output

- Example
  - input  $D_{in} : \vec{x} : [x, y]$
  - output  $D_{out} : \vec{y} :$ 
    - 1 value- what color is is of~ `datatype`
    - 3 values for RGB ~ `uint8`
    - 3 values for the probabilities ~ `P(R), P(G), P(B)`
    - $\vec{y} : [s_r, s_b, s_g]$

$$\vec{y} = \sum_{i=0}^{N-1} \vec{v}_i \phi(\vec{x}_i \cdot \vec{w}_i + b_i)$$

example: say  $N=3$   $N = \#$  of neurons in our 1-layer NN



$w_i = (2,)$   
 $\uparrow$   
 inputs into each neuron

$b_i = \text{scalar}$   
 $\#$  equal to  $\#$  of neurons

$w_i = (2,)$   
 $\uparrow$   
 inputs into each neuron

$b_i = \text{scalar}$   
 .  $\#$  equal to  $\#$  of neurons  
 . no relation to  $I/O$  size

$v_i = (3,)$   
 $\uparrow$   
 match output  
 $[s_r, s_b, s_g]$

- How do we train in batches?

$$\begin{aligned}\vec{x} &\rightarrow (\vec{x})^M \\ (2,) &\rightarrow (M, 2) \\ \vec{y} &\rightarrow (\vec{y})^M \\ (3,) &\rightarrow (3, M)\end{aligned}$$

$$(M, 2) \cdot (2, N) = (M, N)$$

- Loss function

$$Loss_{cross-entropy} = - \sum_{i=0}^{c-1} p_i^{(true)} \log(p_i^{(pred)})$$

- Compares 2 probability distributions (true vs predicted)
- Minimized predictions when  $\vec{p}^{(true)} = \vec{p}^{(pred)}$ 
  - Example:
    - $\vec{p}^{(true)} = [1, 0, 0]; p^{(pred)} = [0.75, 0.2, 0.05]$
    - In the above example, if the loss function encounters such values, it will have the effect of pushing the correct prediction up and the incorrect predictions down to minimize the truth/prediction differences.
- Softmax Cross-Entropy
  - $\vec{s} \rightarrow \text{softmax} \rightarrow \vec{p}$  (scores are given to softmax which are converted into probability-like values)

$$\text{softmax}(\vec{s}) = [\frac{e^{s_0}}{\sum e^{s_j}}, \frac{e^{s_1}}{\sum e^{s_j}}, \frac{e^{s_2}}{\sum e^{s_j}} \dots]$$

where

$$0 \leq \frac{e^{s_0}}{\sum e^{s_j}} \leq 1$$

$$\sum_k (\frac{e^k}{\sum e^{s_j}}) = 1$$

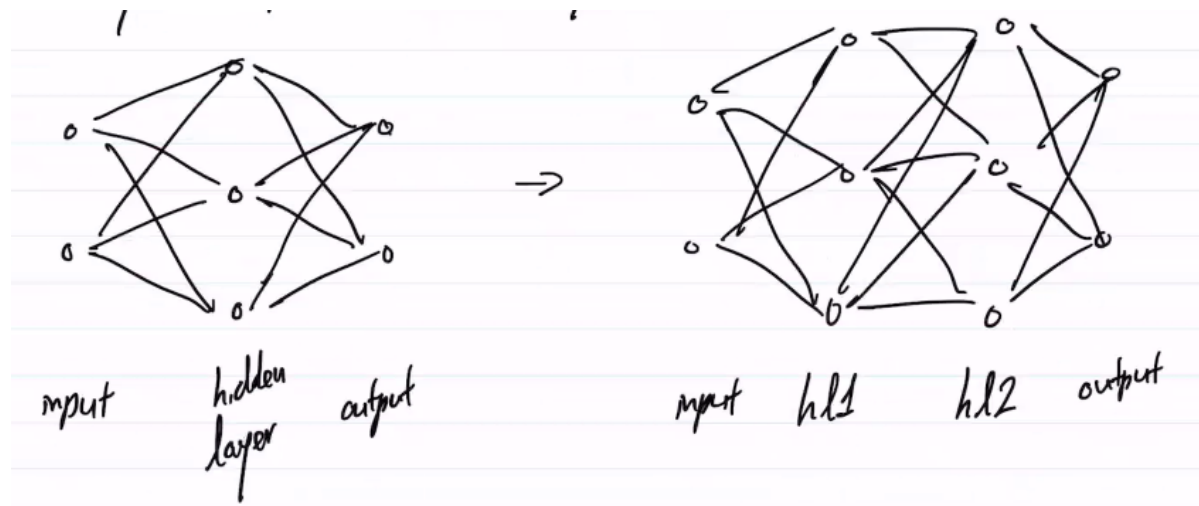
```
import mygrad as mg

# Pseudocode of mg.softmax_crossentropy()
def softmax_crossentropy(scores, labels):
    # scores: (M, D_out)
    # labels: (M)

    # run scores through mg.softmax() to get an array of the same shape
    # transform labels to one-hot encoding (p_true)

    # computes cross-entropy loss between truths, predictions
```

## Multilayer Perceptrons / Multilayer NN



- Connections to the same node from different nodes is equal to the inputting the summation of the results from the different (prior) nodes

- One neuron  $\sim \phi(x_i \cdot w_i + b_i)$

$$L_0 = [\phi(x_0 w_0 + b_0), \phi(x_1 w_1 + b_1) \dots \phi(x_{N-1} w_{N-1} + b_{N-1})] = \phi[XW_0 + b_0]$$

- Previous layer's output are now used as input

$$L_1 = \phi[L_0 W_1 + b_1]$$

$$y = \sum_{i=0}^{N-1} v_i \phi(l_{0,i} w_i + b)$$


- Where  $L$  refers to a whole layer, whereas  $l$  refers to the individual outputs of a previous layer's nodes

## Convolutional Layers and Images

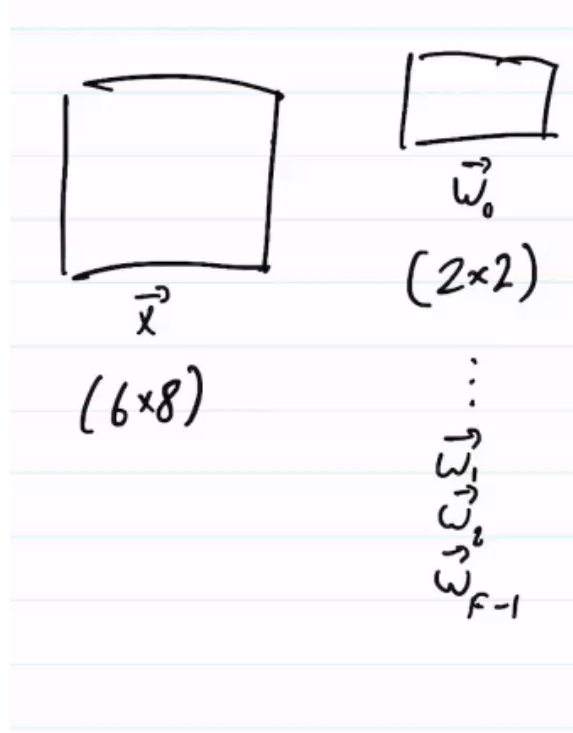


## How can we handle images?


- It is possible to transform an array (representing a matrix of pixels) into a single vector by flattening it and then feeding it to a regular multi-layer neural network (with an output of 10 nodes is using softmax).
  - What are the problems with using images as vectors?
    - The image loses positional information because the main features of the original image are now spread out and disconnected
    - Different neurons have to deal with the same features, depending on where they are in the image
      - Redundant effort
- **Convolution**
  - Filters (sub-matrices with their own weights and biases) “slide” along the image, “looking” at it. *Recall that the dot product of two matrices is equivalent to their similarity. A filter's output from a portion of the image is equivalent to how much the filter “recognizes” that sub-section.*
  - Q: Would a filter activate for a rotation of what it is designed to detect?
    - No, it would not activate optimally. translational invariance  $\neq$  rotational invariance



- Having an image with an image with size (6, 8) and filter size (2, 2) along with stride-1, would yield a new array of size (F, 5, 7) where is F is the number of filters



### ▪ Max Pool



- Operation that often follows one or more convolutional layers
  - Purely a discretization/downsampling process—no learnable parameters
  - $f(\vec{x}) = \max(\vec{x})$
- Why would this be useful?

- high values come from high similarity, therefore max pool helps later parts of the network “focus on” high-similarity detections and reject low-similarity and redundant detections